

HIVE CASE STUDY REPORT

1 Problem Statement

For this assignment, we will be working with a public clickstream dataset of a cosmetics store. Using this dataset, we will extract valuable insights using optimized queries. In addition, we have shown the improvement of the performance after using the optimization techniques on first query.

We have used the data in the link given below:

<https://e-commerce-events-ml.s3.amazonaws.com/2019-Oct.csv>
<https://e-commerce-events-ml.s3.amazonaws.com/2019-Nov.csv>

Both the Dataset, 2019-Oct.csv and 2019-Nov.csv consist of 9 attributes. Description and Data type of all the attributes are given in the below table.

| Attribute_Name | Data_type | Description |
|----------------|-----------|--|
| event_time | timestamp | Time at which the event took place |
| event_type | string | Event type may be 'view', 'cart', 'remove_from_cart', 'purchase' |
| product_id | string | Unique identification of the product |
| category_id | string | Unique identification of the product category. Each product category contains several products |
| category_code | string | Name (if present) of the product category |
| brand | string | Name of the brand |
| price | float | Price of the product |
| user_id | bigint | Permanent user id |
| user_session | string | Identification for the user's session. Remains same for each user's session. It changes every time the user returns back the website after a long pause. |

2 Implementation Phase

2.1 Copying the data set into the HDFS

We have launched an EMR cluster that utilizes the Hive services and connected our EMR cluster with master node. In the next step we have to copy the dataset into the HDFS, therefore, we made the folder 'hive_case_study' under 'tmp' folder.

```
[hadoop@ip-172-31-20-84 ~]$ hadoop fs -mkdir /hive_case_study
[hadoop@ip-172-31-20-84 ~]$ hadoop fs -mkdir /tmp/hive_case_study
[hadoop@ip-172-31-20-84 ~]$ hadoop fs -ls /tmp
Found 4 items
drwxrwxrwt - yarn hdfsadmingroup 0 2021-04-29 04:05 /tmp/entity-file-history
drwxrwxrwx - mapred mapred 0 2021-04-29 03:59 /tmp/hadoop-yarn
drwx-wx-wx - hive hdfsadmingroup 0 2021-04-29 04:06 /tmp/hive
drwxr-xr-x - hadoop hdfsadmingroup 0 2021-04-29 04:16 /tmp/hive_case_study
```

Copy the data from the S3 bucket into the HDFS.

```
[hadoop@ip-172-31-20-84 ~]$ hadoop distcp s3://e-commerce-events-ml/2019-Oct.csv /tmp/hive_case_study/2019-Oct.csv
```

```
[hadoop@ip-172-31-20-84 ~]$ hadoop distcp s3://e-commerce-events-ml/2019-Nov.csv /tmp/hive_case_study/2019-Nov.csv
```

We can observe from the below snapshot that data has been successfully copied to the HDFS.

```
[hadoop@ip-172-31-20-84 ~]$ hadoop fs -ls /tmp/hive_case_study/
Found 2 items
-rw-r--r-- 1 hadoop hdfsadmingroup 545839412 2021-04-29 04:21 /tmp/hive_case_study/2019-Nov.csv
-rw-r--r-- 1 hadoop hdfsadmingroup 482542278 2021-04-29 04:20 /tmp/hive_case_study/2019-Oct.csv
```

2.2 Creating the database and structure of the table

Create database 'clickstream_cosmetics'

```
hive> create database if not exists clickstream_cosmetics ;
OK
Time taken: 0.596 seconds
```

Create the table 'cosmetics' with all the 9 attributes present in the dataset. We have used CSVSerde with the default properties value.

```
hive> create table if not exists cosmetics (event_time timestamp , event_type string , product_id string , category_id string , category_code string , brand string , price float ,user_id bigint , user_session string) row format serde 'org.apache.hadoop.hive.serde2.OpenCSVSerde' with serdeproperties ("separatorChar" = "," , "quoteChar" = "\"", "escapeChar" = "\\") stored as textfile tblproperties ("skip.header.line.count"="1") ;
OK
Time taken: 0.424 seconds
```

Let's look at the 'cosmetics' table schema.

```
[hive> describe cosmetics;
OK
event_time          string          from deserializer
event_type          string          from deserializer
product_id          string          from deserializer
category_id         string          from deserializer
category_code       string          from deserializer
brand               string          from deserializer
price               string          from deserializer
user_id             string          from deserializer
user_session        string          from deserializer
Time taken: 0.039 seconds, Fetched: 9 row(s)
```

Load the data into the table 'cosmetics'

```
[hive> load data inpath '/tmp/hive_case_study' into table cosmetics ;
Loading data to table default.cosmetics
OK
Time taken: 0.591 seconds
```

2.3 Launching Hive queries on the EMR cluster

2.3.1 Find the total revenue generated due to purchases made in October

Query:

```
[hive> select round(sum(price),2) as total_revenue_oct from cosmetics where month(
(to_date(event_time)) = 10 and event_type = "purchase";
```

Output:

```
total_revenue_oct
1211538.43
Time taken: 34.636 seconds, Fetched: 1 row(s)
```

Time taken before optimizing: 34.636 seconds

2.3.1.1 Optimizing the query

We will use optimization technique to improve the hive performance for the above query. Create Partitioned table 'cosmetic_types' based on 'event_type'.

```
[hive> create table if not exists cosmetic_types (event_time string,price string)
partitioned by (event_type string) stored as textfile;
OK
Time taken: 0.085 seconds
```

Let's look at the partitioned table 'cosmetic_types' schema.

```
hive> describe cosmetic_types;
OK
col_name      data_type      comment
event_time    string
price         string
event_type    string

# Partition Information
# col_name      data_type      comment
event_type    string
Time taken: 0.072 seconds, Fetched: 8 row(s)
```

Insert the data into partitioned table 'cosmetic_types'.

```
hive> insert into table cosmetic_types partition(event_type) select event_time, price, event_type from cosmetics;
```

Query on Partitioned Table:

```
hive> select round(sum(price),2) as total_revenue_oct from cosmetic_types where month(to_date(event_time)) = 10 and event_type = "purchase";
```

Output:

```
total_revenue_oct
1211538.43
Time taken: 6.419 seconds, Fetched: 1 row(s)
```

Time taken after Partitioning: 6.419 seconds

We can observe from the above two queries that hive performance is increased after using optimization technique. Time taken to run the query is decreased from 34.636 to 6.419 when we used partitioned table.

2.3.2 Write a query to yield the total sum of purchases per month in a single output.

Query:

```
hive> select month(to_date(event_time)) as month, count(event_type) as total_purchases from cosmetic_types where event_type = "purchase" group by month(to_date(event_time));
```

Output:

```
month    total_purchases
10       245624
11       322417
Time taken: 5.924 seconds, Fetched: 2 row(s)
```

2.3.3 Write a query to find the change in revenue generated due to purchases from October to November

Query:

```
[hive> with total_revenue_month as (select month(to_date(event_time)) as month, sum(price) as total_revenue, lag(sum(price)) over(order by month(to_date(event_time))) as previous_month_revenue from cosmetic_types where event_type = "purchase" group by month(to_date(event_time))) select month, round((total_revenue - previous_month_revenue), 2) as difference_in_revenue from total_revenue_month;
```

Output:

```
month    difference_in_revenue
10       NULL
11       319478.47
Time taken: 6.014 seconds, Fetched: 2 row(s)
```

2.3.4 Find distinct categories of products. Categories with null category code can be ignored

For the above query we will create a table 'cosmetic_prod_category', partitioned based on 'category_code' and store the data into 20 buckets on the basis of 'product_id'.

```
[hive> create table if not exists cosmetic_prod_category (product_id string, category_id string) partitioned by (category_code string) clustered by (product_id) into 20 buckets stored as textfile;
OK
Time taken: 0.099 seconds
```

Let's look at the 'cosmetic_prod_category' table schema.

```
[hive> describe cosmetic_prod_category;
OK
col_name      data_type      comment
product_id    string
category_id   string
category_code  string

# Partition Information
# col_name      data_type      comment
category_code  string
Time taken: 0.192 seconds, Fetched: 8 row(s)
```

Insert the data into table 'cosmetic_prod_category'.

```
[hive> insert into table cosmetic_prod_category partition(category_code) select product_id, category_id, category_code from cosmetics where category_code is not null and category_code <> "";
```

Query:

```
[hive> select distinct(category_code) as categories from cosmetic_prod_category;
```

Output:

```
categories
accessories.bag
apparel.glove
appliances.environment.vacuum
appliances.personal.hair_cutter
furniture.bathroom.bath
furniture.living_room.cabinet
sport.diving
stationery.cartrige
accessories.cosmetic_bag
appliances.environment.air_conditioner
furniture.living_room.chair
Time taken: 10.314 seconds, Fetched: 11 row(s)
```

2.3.5 Find the total number of products available under each category

Query:

```
[hive> select category_code,count(product_id) as total_products from cosmetic_products
category group by category_code;
```

Output:

```
category_code  total_products
accessories.bag 11681
apparel.glove   18232
appliances.environment.vacuum 59761
appliances.personal.hair_cutter 1643
furniture.bathroom.bath 9857
furniture.living_room.cabinet 13439
sport.diving    2
stationery.cartrige 26722
accessories.cosmetic_bag 1248
appliances.environment.air_conditioner 332
furniture.living_room.chair 308
Time taken: 10.776 seconds, Fetched: 11 row(s)
```

2.3.6 Which brand had the maximum sales in October and November combined?

For the above query we will create a table 'cosmetic_brand', partitioned by 'event_type' and bucketing on the basis of 'brand'.

```
[hive> create table if not exists cosmetic_brand (event_time string, brand string,
price string) partitioned by (event_type string) clustered by (brand) into 20
buckets stored as textfile;
OK
Time taken: 0.095 seconds
```

Let's look at the 'cosmetic_brand' table schema.

```

col_name      data_type      comment
event_time    string
brand          string
price          string
event_type     string

# Partition Information
# col_name      data_type      comment

event_type     string
Time taken: 0.046 seconds, Fetched: 9 row(s)

```

Insert the data into 'cosmetic_brand' table.

```

hive> insert into table cosmetic_brand partition(event_type) select event_time,
brand, price, event_type from cosmetics;

```

Query:

```

hive> select brand,round(sum(price),2) as sales from cosmetic_brand where event_
type="purchase" and brand <>" " group by brand order by sales desc limit 1;

```

Output:

```

brand    sales
runail   148297.94
Time taken: 6.837 seconds, Fetched: 1 row(s)

```

2.3.7 Which brands increased their sales from October to November?

Query:

```

hive> with brand_sales as ( select brand,month(to_date(event_time)) as month, ro
und(sum(price),2) as sales, lag(round(sum(price),2)) over(partition by (brand))
as previous_month_sale from cosmetic_brand where event_type ="purchase" and bran
d <>" " group by brand,month(to_date(event_time)) )select brand from brand_sales
where (sales - previous_month_sale) > 0;

```

Output:

| | | |
|---------------|--------------|---------------|
| brand | latinoil | foamie |
| artex | lianail | freedecor |
| balbicare | likato | glysolid |
| batiste | limoni | godefroy |
| beautix | marathon | grattol |
| beautyblender | markell | greymy |
| bioaqua | mavala | igrobeauty |
| bpw.style | metzger | jessnail |
| browxenna | milv | kapous |
| carmex | nagaraku | keen |
| chi | nefertiti | koelf |
| coifin | neoleor | konad |
| concept | nirvel | laboratorium |
| cosima | orly | levissime |
| de.lux | osmo | levrana |
| ecolab | plazan | lovely |
| elskin | polarus | lowence |
| farmona | protokeratin | mane |
| freshbubble | rasyan | marutaka-foot |
| gehwol | refectocil | masura |
| grace | s.care | matrix |
| happyfons | sanoto | maskin |
| haruyama | soleo | missha |
| ingarden | supertan | moyou |
| inn | uno | nitrite |
| insight | uskusi | oniq |
| irisk | yoko | ovale |
| italwax | airnails | profepil |
| jaguar | art-visage | profhenna |
| jas | aura | provoc |
| joico | beauty-free | rosi |
| kaaral | beauugreen | roubloff |
| kamill | benovy | runail |
| kaypro | biore | severina |
| kerasys | blixz | shary |
| kims | bluesky | shik |
| kinetics | bodyton | skinity |
| kiss | candy | skinlite |
| kocostar | cosmoprofi | smart |
| koelcia | cristalinas | solomeya |
| kosmekka | cutrin | sophin |
| lador | deoproce | staleks |
| ladykin | depilflax | strong |
| | dizao | swarovski |
| | domix | tertio |
| | ecocraft | treaclemoon |
| | egomania | trind |
| | elizavecca | veraclara |
| | ellips | vilenta |
| | enjoy | yu-r |
| | entity | zeitun |
| | eos | |
| | estel | |
| | estelare | |
| | f.o.x | |
| | farmavita | |
| | fedua | |
| | finish | |
| | fly | |

Time taken: 12.134 seconds, Fetched: 152 row(s)

152 out of 245 distinct brands increased their sales from October to November.

2.3.8 Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most

For the above query we will create a table 'cosmetic_users', partitioned by 'event_type' and bucketed by 'user_id'.

```

[hive> create table if not exists cosmetic_users (price string, user_id string) p
artitioned by(event_type string) clustered by(user_id) into 10 buckets stored as
textfile;
OK
Time taken: 0.599 seconds

```

Let's look at the 'cosmetic_users' schema.


```
[hive> describe cosmetic_users;
OK
col_name      data_type      comment
price          string
user_id        string
event_type     string

# Partition Information
# col_name      data_type      comment

event_type     string
Time taken: 0.297 seconds, Fetched: 8 row(s)
```

Insert the data into the 'cosmetic_users' table.

```
[hive> insert into table cosmetic_users partition(event_type) select price,user_id,event_type from cosmetics;
```

Query:

```
[hive> with golden_customer_plan as (select user_id, round(sum(price),2) as spend_money, dense_rank () over( order by sum(price) desc) as user_rank from cosmetic_users where event_type = "purchase" group by user_id) select user_id as golden_customer, spend_money from golden_customer_plan where user_rank < 11;
```

Output:

```
golden_customer spend_money
557790271       2715.87
150318419       1645.97
562167663       1352.85
531900924       1329.45
557850743       1295.48
522130011       1185.39
561592095       1109.7
431950134       1097.59
566576008       1056.36
521347209       1040.91
Time taken: 6.23 seconds, Fetched: 10 row(s)
```

2.4 Cleaning up

After completing the hive queries, we have dropped the database and terminated the cluster.

```
[hive> drop database clickstream_cosmetics;
OK
Time taken: 0.055 seconds
```